

07 · Did the regional campaign work? — synthetic control ANCHOR B

July 9, 2026

1 07 · Did the regional campaign work? — synthetic control ANCHOR B

The business decision. We ran a TV campaign in one metro area — a **DMA** (*Designated Market Area*, the standard US media-buying region) — and left the rest of the country dark. Sales in that metro are up since launch. Before we spend to roll the campaign out nationally: **was the lift caused by the campaign, or would sales have risen there anyway?**

1.0.1 Why this is hard: there was no control group

Unlike the uplift notebook (01), here there is **no randomized hold-out** — we didn’t flip a coin to decide which metro got the campaign, we just picked one. So we can’t simply compare it to “the others.” And we can’t just compare the treated metro’s sales *after* vs *before*, because sales everywhere are pushed around by a shared **trend** (secular growth), **seasonality**, and a common **macro wave** (the economy, category demand). A naive before/after would credit the campaign with an uplift the market would have delivered on its own. We need the one thing we can never observe: **what sales in that metro would have been if we had not run the campaign** — its *counterfactual*.

1.0.2 The idea: build the counterfactual from the other markets

Synthetic control (Abadie & Gardeazabal, 2003) constructs that missing counterfactual as a **weighted blend of the untreated markets** — the “**donor pool**.” It searches for a set of weights so that the blend closely tracks the treated market’s sales *in the pre-launch period*. If a weighted combination of, say, 60% Denver + 25% Phoenix + 15% others matched our metro week-for-week *before* the campaign, then — assuming those markets keep moving together — that same blend is our best guess of what our metro *would* have done *after*, absent the campaign. The **gap** between the metro’s actual sales and this “synthetic” version, after launch, is the estimated lift.

The reason it can work at all: all markets share the same underlying **factors** (trend, season, macro). A blend of donors that matches the treated unit’s pre-period is, in effect, matching it on those shared factors — so the blend and the treated unit stay locked together *unless something (the campaign) breaks them apart*.

We do this the Bayesian way, so instead of one number we get a **posterior** — a full distribution of plausible lifts — which gives us honest uncertainty, a real significance test (a **placebo per-**

mutation, defined in Depth B), and robustness checks a VP of Growth can interrogate before committing national budget.

1.0.3 How this notebook is organised

The **7-step contract** (question · simulate a known truth · identify · estimate · validate · decide in € · caveats) plus the full **three depths**:

- **Depth A** — why the tempting naive estimators fail (a bake-off)
- **Depth B** — placebo inference in *space* and *time*, with the scale-free Abadie RMSE-ratio test
- **Depth C** — the euro rollout decision and its sensitivity to assumptions

FAST=False SC sampling={'draws': 1500, 'tune': 1500, 'chains': 4}


1.1 2 · Simulate a ground truth (why fake data first)

As in every notebook here, we **cannot grade a causal method on real data** — the counterfactual (“what our metro would have done without the campaign”) is never observed. So we build a simulator where we *plant* the true lift, confirm the method recovers it, and only then trust it.

The simulator produces a weekly **sales panel**: a table with one row per market per week, for 30 DMAs over 60 weeks. Each market’s sales are built from **shared latent factors** — a rising trend, a seasonal cycle, and a random macro wave common to all markets — combined with that market’s own baseline level, its own sensitivity to each factor (“factor loadings”), and idiosyncratic noise. The campaign launches in **week 40** in market 0 only, adding a **true +12% lift** from then on.

The shared macro wave is the villain: it moves *every* market together, so it’s exactly what a naive before/after would mistake for campaign effect. But it’s also the hero of the method — because the factors are *shared*, a weighted blend of the *other* markets can reconstruct market 0’s untreated path. That is precisely what makes the effect *identifiable* here (recoverable in principle), and the plot below shows the setup: one dark line (treated) inside a bundle of grey donors, all riding the same wave.

Treated: dma_00, launch week 40. TRUE lift €17.0k/week, total €339k.



1.2 3 · Identify — the estimand and the assumptions it rests on

The estimand (the precise thing we *want*). For each post-launch week t , the causal effect on the treated market is the potential-outcome contrast

$$\tau_t = Y_{1,t}(1) - Y_{1,t}(0), \quad t > T_0,$$

where $Y_{1,t}(1)$ is **observed** (the campaign ran) but $Y_{1,t}(0)$ — what week- t sales *would have been* with no campaign — is the **missing counterfactual**.

The estimator. Synthetic control fills in that counterfactual with a weighted average of donor markets, so the effect we actually compute is

$$\hat{\tau}_t = Y_{1,t} - \underbrace{\sum_j w_j Y_{j,t}}_{\widehat{Y_{1,t}(0)} = \text{synthetic control}}, \quad t > T_0,$$

where $Y_{1,t}$ is the **treated** market’s sales in week t , $Y_{j,t}$ is **donor** market j ’s sales, T_0 is the launch week, and w_j are the **donor weights**. The weights are chosen to match the treated market’s *pre-launch* path and are constrained to the **simplex**: each $w_j \geq 0$ and they sum to 1 ($\sum_j w_j = 1$). Think of the simplex constraint as “the synthetic market must be a genuine *weighted average* of real markets” — no negative weights, no scaling beyond the observed range.

Why the simplex matters (the core of Abadie’s method). If instead we let weights be any numbers (ordinary regression), the “synthetic” market could sit *outside* the range of any real market — **extrapolation** dressed up as a match, which fits the pre-period beautifully and then diverges wildly after. Forcing weights onto the simplex keeps the synthetic inside the **convex hull** of the donors (loosely, “inside the cloud of real markets”), which is what makes the post-launch projection credible.

The assumptions — each gets an explicit check later:

assumption	plain meaning	checked in
convex hull / no extrapolation	the treated market lies inside the donor cloud; simplex weights enforce it	weights (Step 4)
good pre-fit	the synthetic actually tracks the treated market <i>before</i> launch	pre-RMSE gate (Step 5)
no anticipation	nothing shifts sales <i>before</i> launch (no leaked campaign, no pre-buying)	placebo-in-time (Depth B)
stable factor structure	donors and treated keep co-moving <i>after</i> launch (the loadings matched pre-launch still hold)	long pre-period + placebo-in-time

assumption	plain meaning	checked in
no spillover / interference	the campaign in market 0 doesn't change donor markets (SUTVA for geos)	assumed by design (simulation); on real data, exclude adjacent / media-overlapping donors and re-run — placebo-in-space <i>cannot</i> detect spillover (it contaminates every placebo too)
no concurrent treated shock	nothing else hit market 0 exactly at launch	untestable — institutional knowledge only


1.3 4 · Estimate — a Bayesian synthetic control

We put a **Dirichlet prior** on the donor weights and fit them to the pre-period. (The *Dirichlet* is the natural probability distribution over the simplex — draws from it are automatically non-negative and sum to 1 — so it's the Bayesian way to say “the weights live on the simplex.”) Fitting returns a **posterior** over the weights, and therefore a posterior over the whole counterfactual path and the lift — Abadie's constrained least squares, but with uncertainty for free.

Two things to read off the fit: - **Which donors** the synthetic leans on (the weight bars). One caveat to state plainly: the Dirichlet prior produces a *regularized blend* spread over many donors, rather than classical synthetic control's sparse handful. That's **more stable** out-of-sample but **less individually interpretable**. We summarise the spread as an “**effective number of donors**” $= 1 / \sum_j w_j^2$ (the inverse Herfindahl index — 1 if all weight is on one donor, J if spread evenly over J donors). - **Pre-fit RMSE** — the root-mean-square gap between the treated market and its synthetic *before* launch, in €000. Small means the synthetic is a good stand-in; this is the precondition for believing the post-launch gap — Step 5 turns this into a concrete PASS/FAIL gate (pre-RMSE below ~a third of the weekly lift it must detect).

On real data. You don't need our simulator — this method runs on *your own* weekly (or daily) sales-by-market panel, which every company already has: one treated market, the rest as donors. The canonical *public* example is Abadie's **California Proposition 99** tobacco study (California vs 38 other US states); it has the identical shape (one treated unit, many donors, a known intervention date) and is the standard dataset to reproduce this on.

SC sampling convergence: max r-hat 1.000 - min ESS 3994 - divergences 0
pre-fit RMSE €2.86k · estimated total lift €305k (true €339k) · 90% CI
[€270k, €339k]




1.4 5 · Validate — recovery, calibration, the pre-fit gate, and robustness

Four checks — and one of them delivers a genuinely humbling result we keep rather than hide:

1. **Recovery & calibration.** The true per-week path should sit inside the posterior band. That band is a **posterior predictive**: it carries not only *which-donors* (weight) uncertainty but the treated market’s own idiosyncratic noise. That noise term is what makes the per-week band near-nominal; a weights-only band is too narrow and can leave the truth *outside* a “90%” interval.
2. **Multi-seed calibration.** One panel is a single draw of the world, so we refit on fresh panels and report the recovery bias and coverage. The honest result: the **per-week** band is near-nominal, but the **total-lift** interval covers well *below* nominal — synthetic control’s genuine sample-to-sample variance is larger than any single panel’s posterior width. So the posterior interval is an *optimistic* guide to the total, and the trustworthy inference is the **placebo permutation test** (design-based, calibrated by construction), not the band. This is *why* the euro call below is “confirm with a second test,” not “the interval says go.”
3. **Pre-fit gate** — the pre-period gap must hover around **zero**: concretely, pre-RMSE below about a third of the weekly lift the design must detect (printed as an explicit PASS/FAIL below). A bad synthetic match makes the post gap meaningless — this is the precondition, not a formality.
4. **Leave-one-out donor robustness** — drop each of the top-weighted donors and refit; if the estimate lurches when one donor leaves, the result hangs on a single market and is fragile.

pre-fit gate: pre-RMSE €2.86k vs $|weekly\ lift|/3 = €5.1k \rightarrow PASS$
 pre-launch gap mean €0.30k (~ 0 = good fit). L00 total range €309-337k vs the same-fitter SLSQP baseline €318k (Bayesian €305k) - no single donor drives it.



Across 24 fresh panels - recovery bias €+23k (+7%), with a LARGE seed-to-seed sd of €59k.

PER-WEEK band covers the true path 86% of the time - ~ nominal; the posterior-predictive obs-noise term earns this.


But the TOTAL-lift 90% interval covers only ~50% across panels: SC's sample-to-sample variance dwarfs any single panel's posterior width, so the total interval is optimistic.

=> the CALIBRATED inference is the placebo permutation test (below), not the posterior band alone.

1.5 6 · Decide, in euros

The campaign cost **€300k**. The rollout question is not “is the lift positive?” but “**what’s the probability the lift beats what we paid?**” — read off the posterior of total incremental sales, plus how much headroom before the economics turn.

```
{
  "P_value_gt_cost": 0.5878333333333333,
  "expected_value": 304.7713414404065,
  "expected_net_value": 4.771341440406502,
  "expected_roi": 0.015904471468021673,
  "value_lo": 270.2408765084909,
  "value_hi": 339.44759544458793,
  "cost": 300.0,
  "decision": "TEST FURTHER"
}
```




Value of information: resolving the cost-line uncertainty is worth ~€6k of expected decision value.

A confirmatory geo test averaging ~33 independent treated markets would tighten the interval

enough to reach $P(\text{lift} > \text{cost}) \geq 0.90$ - and since the multi-seed check showed this interval is

optimistic, treat 33 as a LOWER bound on the test size. That is the experiment to run, not a blind launch.



1.6 7 · Caveats

- **Here synthetic control is high-variance, not conservative.** In this panel the point estimate lands ~10% low, but across seeds (Step 5) it is roughly *unbiased* with large sample-to-sample spread — and its posterior interval is *optimistic* about the total lift. So trust the **placebo permutation test**, not the point estimate or the band.
- **Pre-fit is a gate, not a formality.** High pre-RMSE => the post gap is uninterpretable;

no Bayesian machinery fixes a bad donor pool.

- **No-anticipation & no-spillover are assumptions**, stress-tested by the placebos in Depth B.
- **One treated unit = limited power**. Significance comes from the permutation test, coarse with few donors.


2 Depth A · Why the naive estimators fail

Two tempting shortcuts, both confounded by the shared trend/season/macro wave:

- **Before/after**: post mean — pre mean of the treated DMA — books the entire trend as effect.
- **Treated — average control**: the average control has the wrong factor loadings, so it doesn't cancel the common wave cleanly.

Synthetic control is the fix: a *weighted* control chosen to match pre-launch.

```
{  
  "true": 17.0,  
  "before/after": 24.3,  
  "treated \u2212 avg ctrl": 27.1,  
  "DiD vs avg ctrl": 14.5,  
  "synthetic control": 15.2  
}
```



Read-out. Before/after badly overstates the lift (it books the trend); treated-minus-average is **worse** — it books the raw between-market level gap, differencing nothing out; differencing out the shared trend (DiD vs the average control) removes the trend bias yet a *level* gap from mismatched factor loadings remains; **only the reweighted synthetic control lands on the truth**. Same data, three answers — the estimator *is* the identification strategy.

3 Depth B · Placebo inference — is the gap statistically real?

We have a lift estimate, but with **one treated market** there's no classical standard error to test it (standard errors assume many independent units). Instead we ask a **permutation** question, the inferential heart of synthetic control:

If the campaign had no effect, how unusual would a post-launch gap this big be, just by chance?

Placebo-in-space. We answer it by pretending, in turn, that *each donor market* was the treated one: fit that donor its own synthetic control from the remaining donors, and record its post-launch “gap.” Since none of the donors actually got a campaign, these placebo gaps map out the distribution of gaps we’d see from noise alone. If the *real* treated market’s gap sits far out in the tail of that placebo distribution, a fluke is unlikely. Two refinements from Abadie:

- **Discard placebos with a poor pre-fit** — a donor the synthetic can’t match before launch produces a meaningless gap that would pollute the comparison.
- **Use the RMSE-ratio statistic**, not the raw gap: the *post*-period RMSE divided by the *pre*-period RMSE. This is **scale-free** (fair across big and small markets) and, crucially, doesn’t reward a placebo that simply fits badly *everywhere* (high post *and* high pre $\rightarrow$ ratio ~ 1).

The **p-value** is then the treated market’s rank in the ratio distribution: the fraction of markets (placebos + treated) whose ratio is at least as extreme as the treated market’s. A small p — say ≤ 0.05 — means “only this market, out of all of them, shows a jump this sharp relative to its own pre-fit.”

placebo-in-space kept 18/29 donors (dropped 11: pre-RMSE > 5x the treated fit)

placebo p vs filter cutoff: 2x $\rightarrow$ p=0.056 · 5x $\rightarrow$ p=0.053 · 20x $\rightarrow$ p=0.333


(tighter filters compare only well-matched placebos and sharpen the test; a

loose 20x filter admits poor synthetics and dilutes it)

placebo p (gap) = 0.053 [add-one permutation, treated excluded from placebos] ·

RMSE-ratio p = 0.033 [Abadie rank/N, treated included] (treated ratio 6.7 vs


placebo median 1.3; both rank the treated 1st)



Placebo-in-time. A second falsification: pretend the campaign launched *before* it did (a fake

week inside the pre-period). The estimated “effect” should be ~ 0 — a non-zero fake effect would mean the method invents lift (anticipation or overfitting).


placebo-in-time avg gap €0.9k vs pre-fit RMSE €2.86k (should be ~ 0 , and it is well within the pre-fit noise) - no spurious pre-effect.



4 Depth C · The euro rollout decision & sensitivity

Pull it together into the number a VP signs off on, and pressure-test how the call moves with the **launch-timing** assumption (did we mis-date the campaign?) and the **donor pool**.

```
{
  "true_total": 339.4806333100736,
  "estimated_total": 304.7713414404065,
  "ci90": [
    270.2408765084909,
    339.44759544458793
  ],
  "pre_rmse": 2.855007246428569,
  "p_space": 0.05263157894736842,
  "p_rmse_ratio": 0.03333333333333333,
  "placebo_in_time": 0.9046168542429143,
  "P_lift_gt_cost": 0.5878333333333333,
  "decision": "TEST FURTHER",
  "eff_donors": 14.579714199645526
}
```



4.0.1 The one-paragraph decision

The campaign drove a **real** incremental lift: the synthetic tracks the treated market closely before launch (low pre-RMSE, pre-gap ~ 0), spreads its weight across ~ 15 effective donors (no single market dominates), and it is robust to dropping any single donor (leave-one-out stays close to the same-fitter baseline). Nudging the *assumed* launch date does shift the estimate — but that is mechanical, not fragility: mis-dating folds no-effect weeks into the post window and dilutes the average lift, so the peak at the true launch week is itself a diagnostic. The post gap is an outlier on both the raw-gap and the scale-free RMSE-ratio **placebo permutation** tests, and the time placebo shows no spurious pre-effect — that permutation test is the calibrated evidence the effect is real. The euro call is deliberately cautious: the point estimate carries meaningful seed-to-seed variance (this panel lands $\sim 10\%$ below truth; across panels the estimator is roughly unbiased but *high-variance*), and the multi-seed check shows the **total-lift interval is optimistic** — so we do **not** let a single posterior interval authorise a national rollout. Against the €300k spend the effect sits near break-even: “**promising, confirm with a second geo test,**” **not** “**scale nationally today.**” The effect is clearly real; whether it clears *this* cost is the close call — and one an optimistic interval should make us *more*, not less, cautious about.