

08 · Did the loyalty rollout work? — difference-in-differences (CausalPy)

July 13, 2026

Runs in the legacy environment (pymc<6): `make env-legacy, kernel cmp-legacy.`

The business decision. We launched the loyalty app in some stores, not others. Revenue in the launch stores went up — but revenue went up *everywhere* (a seasonal tide lifts all boats). Did the app add anything **net of that tide**, and is it worth rolling out to the whole chain?

The idea: two differences

Difference-in-differences (DiD) is the workhorse of quasi-experiments. It takes **two** differences and subtracts them:

$$\text{DiD} = \underbrace{(\bar{Y}_{\text{after}}^{\text{treated}} - \bar{Y}_{\text{before}}^{\text{treated}})}_{\text{treated change}} - \underbrace{(\bar{Y}_{\text{after}}^{\text{control}} - \bar{Y}_{\text{before}}^{\text{control}})}_{\text{control change}}.$$

The first difference (treated after — before) contains the app effect **plus** the seasonal tide; the second (control after — before) is the tide **alone**. Subtracting removes the tide and leaves the app effect. It’s the same “compare like with like” logic as everywhere else, applied *across time*.

The one assumption everything rests on: parallel trends

DiD is valid only if, **absent the app**, treated and control stores would have moved *in parallel* — the control’s change is a fair stand-in for what the treated stores’ change would have been. This is untestable after launch (we never see the treated stores’ no-app future), but the **pre-launch trends are the evidence**: if the two groups moved together *before* the app, parallel trends is credible. We check this with an **event study** (the effect period-by-period) and falsify with a **placebo** (a fake launch date in the pre-period, which should show ≈ 0 effect).

On real data. Swap in your **own store/region panel** — one row per store per period with revenue, a treated/control flag, and a period index. The textbook public example is **Card & Krueger (1994)** on minimum wage and employment (fast-food restaurants in New Jersey vs Pennsylvania). One warning we *demonstrate live* in Step 7: if stores adopt at **different times** (“staggered rollout”), naive DiD can be badly biased and you need modern estimators (Callaway–Sant’Anna, Sun–Abraham).

It follows the repo’s fixed **7-step contract**: question → simulate → identify → estimate → validate → decide in € → caveats.

2 · Simulate a ground truth

40 stores, half get the loyalty app at week 12. Every store shares a seasonal pattern and has its own baseline; the app adds a **true €400/store/week**. Treated and control move in **parallel** before launch by construction — so DiD should recover €400, and the event study should show a flat pre-trend.

The data-generating model — exactly what `dgp.did_rollout` implements (defaults & seed in `src/cmp/dgp.py`). Stores $s = 1, \dots, 40$ (a random half treated, $D_s = 1$), weeks $t = 0, \dots, 23$, launch at week 12:

$$Y_{st} = \underbrace{\alpha_s}_{\text{store baseline}} + \underbrace{60 \sin\left(\frac{2\pi t}{12}\right)}_{\text{shared season}} + \underbrace{400 D_s \mathbf{1}[t \geq 12]}_{\text{true effect}} + \varepsilon_{st}, \quad \alpha_s \sim \mathcal{N}(1000, 150^2), \quad \varepsilon_{st} \sim \mathcal{N}(0, 80^2).$$

Parallel trends holds by construction: the seasonal term is *identical* across stores (no store-specific loading on it) and α_s is time-invariant, so absent the app the treated and control means move in lockstep — exactly what the event study’s flat pre-trend should (and does) show. This is the assumption the caveats section then breaks on purpose (staggered adoption) to demonstrate the two-way fixed-effects (TWFE) trap.

TRUE effect = €400/store/week * 40 stores, 24 weeks, launch week 12

	store	unit	t	week	group	post_treatment	post	revenue
0	store_00	0	0	0	1	False	0	1071.613014
1	store_00	0	1	1	1	False	0	1066.597066
2	store_00	0	2	2	1	False	0	1019.187212
3	store_00	0	3	3	1	False	0	834.246887
4	store_00	0	4	4	1	False	0	1030.278745

3 · Identify — the estimand, and the two assumptions that buy it

The estimand, precisely — the *estimand* is the exact quantity we are trying to estimate. Store s in week t has two *potential* revenues: $Y_{st}(1)$ with the app and $Y_{st}(0)$ without — we only ever observe one of them. With $D_s \in \{0, 1\}$ flagging the pilot stores and t_0 the launch week, DiD targets the **ATT**, the average treatment effect *on the treated stores, after launch*:

$$\text{ATT} = \mathbb{E}[Y_{st}(1) - Y_{st}(0) \mid D_s = 1, t \geq t_0].$$

The second term inside is the counterfactual we never see — what the pilot stores *would have* earned post-launch without the app. Note the final **T**: this is the effect *on the pilot stores*, not on the whole chain; it need not equal what the other stores would experience (pilot sites are rarely picked at random), which is exactly why Step 6 stress-tests the transfer before scaling to 500 stores. Two named assumptions let the control stores stand in for the missing counterfactual:

Assumption 1 — Parallel trends (PT), stated on *untreated* potential outcomes: for any post-launch week t and pre-launch week t' ,

$$\mathbb{E}[Y_{st}(0) - Y_{st'}(0) \mid D_s=1] = \mathbb{E}[Y_{st}(0) - Y_{st'}(0) \mid D_s=0], \quad t \geq t_0 > t'.$$

In words: had the app never launched, the pilot stores’ revenue would have *changed* by the same amount as the controls’. **Levels may differ freely** — pilot stores can be systematically bigger, busier, richer; it is the *changes* that must match. In marketing terms: the seasonal tide, chain-wide promotions and macro shocks hit both groups alike, and nothing *time-varying and group-specific* (a regional campaign running only in pilot cities, say) is in play.

Assumption 2 — No anticipation (NA):

$$\mathbb{E}[Y_{st}(1) - Y_{st}(0) \mid D_s = 1] = 0 \quad \text{for all } t < t_0,$$

i.e. the app moves nothing before it exists. Sounds vacuous; isn’t. Pre-launch buzz marketing, staff trained weeks early, or customers *deferring* purchases until the app’s perks arrive all violate it — and they contaminate the very “before” baseline that DiD subtracts. Step 7 breaks this assumption on purpose so you can see what the damage looks like.

Why DiD = ATT — the one-line derivation. Take the treated group’s observed before→after change and split it, using NA (the treated pre-period observation *is* $Y(0)$), into effect plus counterfactual trend (throughout, superscript 1 = the treated/pilot group mean, 0 = the control group mean):

$$\bar{Y}_{\text{post}}^1 - \bar{Y}_{\text{pre}}^1 = \underbrace{\text{ATT}}_{\text{what we want}} + \underbrace{\mathbb{E}[Y_{st}(0) - Y_{st'}(0) \mid D_s=1]}_{\text{counterfactual trend}}.$$

PT says that trend equals the *control* group’s observed change $\bar{Y}_{\text{post}}^0 - \bar{Y}_{\text{pre}}^0$. Subtract it:

$$\text{DiD} = (\bar{Y}_{\text{post}}^1 - \bar{Y}_{\text{pre}}^1) - (\bar{Y}_{\text{post}}^0 - \bar{Y}_{\text{pre}}^0) = \text{ATT}.$$

That is the entire method: two assumptions in, one subtraction out.

The regression forms you’ll meet. The **2×2 DiD** (two groups — pilot vs control — by two periods — pre vs post) is the interaction coefficient β_3 in $Y = \beta_0 + \beta_1 \text{group} + \beta_2 \text{post} + \beta_3 (\text{group} \times \text{post}) + \varepsilon$ — the version we fit in Step 4. Its dynamic cousin, the **event-study regression**, estimates one effect per week relative to launch:


$$Y_{st} = \alpha_s + \lambda_t + \sum_{k \neq -1} \beta_k \mathbf{1}[t - t_0 = k] D_s + \varepsilon_{st},$$

with store fixed effects α_s , common week effects λ_t , normalized at $k = -1$ — i.e. we peg the last pre-launch week to zero and read every other week’s effect *relative to it* (one week must be dropped as the reference, or the store and week effects would be collinear with the full set of β_k). That is why the plot’s dot at $k = -1$ sits exactly on zero and every other dot is a gap *from* it. The nonparametric per-week treated–control gap we plot below *is* this regression’s β_k in the balanced two-group case — same picture, so you can map our plot onto any paper’s event study. The **leads** ($k < 0$) test NA directly and make PT *credible*; nothing can test PT itself post-launch (we never observe the treated stores’ no-app future).

What the DAG buys us. Contrast with notebook 01: backdoor adjustment needs the confounder *measured* so you can condition on it. DiD strikes a different bargain — the store-level


confounder α_s (location, size, clientele: whatever drives both “chosen as pilot” *and* revenue) may stay **unobserved**, provided it is *time-invariant*: differencing each group against its own past removes α_s , and subtracting the control change removes the common shocks λ_t . What survives both subtractions is precisely what PT rules out: time-varying, group-specific shocks.

DiD's bargain: the confounder α_s can stay unobserved — if time-invariant



within-group differencing removes α_s ; subtracting the control change removes λ_t

Pre-launch leads scatter around 0 within their 90% bands (0 of 12 exclude 0):
 mean |gap| €22, on the order of the per-week noise SE €57 -
 consistent with parallel trends, not a pre-trend.
 (Centred leads average exactly 0 by construction, so it is their SPREAD vs
 the noise band, not a mean, that carries the evidence.)
 Post-launch gaps jump to ~€400.



How to read the event study (right panel). This is the single most important plot in a DiD analysis. Each dot is the treated-minus-control revenue gap in one week, relative to launch (week 0). Each dot now carries a **90% uncertainty band**. The story we *want* to see, and do: the pre-launch dots are **flat and their bands comfortably include zero** — that is the honest

content of “parallel trends look credible” (a lead whose band *excluded* zero would be the red flag), and the assumption the whole method rests on — and then they **jump to $\approx \text{€400 at launch}$** and stay there. A rising or falling pre-launch trend would have been a red flag that the two groups were already diverging, and we’d have had to abandon DiD for synthetic control (notebook 07). The left panel shows the raw series so you can see both groups riding the same seasonal wave.

The same picture, twice — and why we redraw it in a moment. The gaps above are computed *nonparametrically*: a per-week treated-minus-control difference of means, with a band built from the between-store spread in that week alone. Step 0 immediately below re-derives exactly these β_k from a **regression** with store and week fixed effects, and gives them **store-clustered** standard errors — the bands that a referee would ask for. Same dots, honest error bars; compare the two panels as you go.

Step 0 · The classical read (no likelihood, no priors, no sampler)

Before any Bayesian machinery, ask what a competent analyst would do here **without** it. The answer is not a different analysis: it is the *same estimand* (the ATT of §3), identified by the *same two assumptions* (parallel trends + no anticipation), estimated with the simplest tool that is still correct. The causal work lives in the identification argument, not in the machinery — and it is worth seeing that once, cleanly, before the likelihood and the priors arrive and start doing work of their own.

The estimator, naked. The canonical **2×2 DiD** is four cell means and two subtractions:

$$\widehat{\text{ATT}}_{2 \times 2} = (\bar{y}_{\text{post}}^T - \bar{y}_{\text{pre}}^T) - (\bar{y}_{\text{post}}^C - \bar{y}_{\text{pre}}^C),$$

with T the pilot stores and C the controls. Run as a regression on the store-week panel,

$$Y_{st} = \beta_0 + \beta_1 D_s + \beta_2 \text{post}_t + \beta_3 (D_s \times \text{post}_t) + \varepsilon_{st},$$

the interaction coefficient β_3 **is** that difference of differences, algebraically — no approximation. So `cl.did_2x2` returns one number, and it is the ATT.

The covariance choice is not a detail — it is the estimate’s other half. A point estimate without a defensible standard error is half an answer, and the *default* standard error here is indefensible. Our 960 rows are not 960 independent observations: they are 40 stores watched for 24 weeks, and a store’s weeks resemble each other (a shared baseline α_s ; in real chains, serially-correlated local shocks on top). The information content of the panel is closer to **40** than to 960. **Clustering on store** is the classical fix — it lets the errors be arbitrarily correlated *within* a store and asks only for independence *across* stores, so the 40 stores, not the 960 rows, become the unit of inference. That correction is the point of **Bertrand–Duflo–Mullainathan (2004)**, whose famous finding is that ignoring within-panel correlation can leave DiD standard errors off by a factor of five or more.

We therefore fit the same regression **twice** — once with iid (“nonrobust”) errors, once clustered on store — and print both. The point estimate will be identical to the last decimal; only the story about its precision changes. Watch which direction it moves, because it is not the direction most people expect.

THE ESTIMAND - the ATT, i.e. the 2x2 DiD interaction (960 store-weeks, 40 stores).

One point estimate; two stories about how precise it is:

```
iid      : 390.49€ [90% CI 350.95€, 430.03€] * SE 24.02 * nonrobust
clustered : 390.49€ [90% CI 370.06€, 410.92€] * SE 12.41 * cluster-robust
by unit (40 clusters)
-> clustering multiplies the DiD's SE by x0.52
```

A DIFFERENT contrast, from the very same fit - the pilot-vs-control LEVEL gap:
iid SE €17.0 vs clustered SE €51.5 -> the iid SE was x3.0 TOO NARROW.

GRADE vs the planted truth (€400/store/week): the estimate is off by €-10,
and the truth falls INSIDE the clustered 90% CI.

BDM's mechanism, isolated: the SAME panel with serially-correlated (AR(1)) weekly shocks.

rho	mean clustered SE	mean iid SE	ratio	(10 panels per rho)
0.00	9.4	23.6	0.40	
0.50	14.9	23.5	0.63	
0.80	20.4	23.4	0.87	
0.95	18.1	23.2	0.78	

-> as persistence rises, the DiD's TRUE (clustered) SE roughly doubles, while the iid SE sits there unmoved: the nonrobust formula cannot see serial correlation at all. That blindness is the gap BDM measured.
(The ratio dips again at rho=0.95: a near-permanent shock behaves like a store baseline and cancels in the DiD contrast. It is INTERMEDIATE persistence - outliving the pre/post boundary, but not the panel - that does the most damage.)

Read the two ratios together — the lesson is not “clustering widens intervals”. On our estimand (the interaction) clustering *shrank* the standard error, roughly halving it. On the level gap from the identical regression, clustering *inflated* it several-fold. Both are correct, and the mechanism is the same object seen from two sides: each store carries a time-invariant baseline α_s (sd €150 by construction — larger than the €80 weekly noise).

- The **level gap** compares *different stores*, so it is exposed to that baseline scatter — but the iid formula, believing it has 960 independent rows, divides the scatter by $\sqrt{960}$ instead of by $\sqrt{40}$. It reports a precision the data never had. This is the classic anti-conservative failure, and it is the one that ends careers: a confident interval around a number you cannot actually resolve.
- The **DiD interaction** compares each store *against its own past*, so α_s cancels out of the contrast entirely. But it does *not* cancel out of the pooled residual: iid OLS still charges the estimate for a baseline scatter it never used. The cluster-robust sandwich looks at the 40 store-level trajectories directly, notices the cancellation, and hands back the honest — here, *smaller* — standard error.

So the iid SE is not “conservative” or “anti-conservative” as a rule. It is simply **wrong in a direction you cannot predict without doing the clustering**, and the sign of the error flips between two coefficients of the same fit. That is the whole argument for making the covariance choice explicit rather than accepting a default.

The rho sweep then isolates **Bertrand–Duflo–Mullainathan’s own case**: keep everything, and make only the weekly shocks persistent (a good month tends to follow a good month, as real store revenue does). Read the two SE columns down the table. The **clustered** SE climbs steeply with persistence — that is real uncertainty, and it is really there. The **iid** SE barely twitches, because serial correlation is not a quantity that appears anywhere in its formula: the nonrobust estimator is not *underestimating* this uncertainty so much as it is structurally unable to perceive it. Persistence could double the true sampling variability of your DiD and the default **statsmodels** interval would report the same width it always did, with the same confident three decimal places. Our simulated chain is unusually kind on this axis (its weekly shocks are iid by construction); **your** panel will not be, which is why clustering on the panel unit is not optional.

Hold on to this number. The clustered standard error above is the classical read’s honest width, and it will come back — pointedly — in **5x**, where we lay it next to the Bayesian posterior.

PRE-PERIOD (11 leads, $k < 0$) - the parallel-trends evidence:


0 of 11 lead CIs exclude zero (at 90%, ~1 in 10 would by chance alone).
mean |lead| €23, against a typical lead SE of €36 - noise, not a pre-trend.

POST-PERIOD (12 lags, $k \geq 0$) - the dynamic ATT:

mean €384 (true €400), range €349–€441: an immediate, flat step - no build-up, no decay.

averaging the lags reproduces the 2x2 ATT above (€390) to within €6.

What this 90% confidence interval does NOT say: that there is a 90% probability the true effect lies inside it. It is a property of the *procedure* - 90% of intervals built this way would cover the truth across repeated samples. This interval either contains the truth or it does not. To get a probability *about the effect itself* - the thing a go/no-go rule like $P(\text{lift} > \text{cost}) \geq 0.9$ actually needs - you need a posterior. That is what the Bayesian section adds.



Read-out — the classical answer, in business terms. The loyalty app lifted revenue in the

pilot stores by the amount printed above, per store per week, net of the seasonal tide; the 90% clustered confidence interval is narrow enough that the entire interval sits far above any plausible running cost for the app (§6 does that arithmetic properly). The planted truth is €400 and it sits inside the interval — the classical estimator is doing its job. The event study says the same thing with more resolution: **eleven flat leads** (there are twelve pre-launch weeks, but one — $k = -1$ — is spent as the reference and pinned to zero by construction, so eleven are free to move and none of them does), then an immediate step to $\approx$ €400 that holds for twelve weeks. The leads are the evidence that makes parallel trends credible — and the *only* evidence we will ever have for it. If the leads had trended, this notebook would have ended here, and we would be reaching for synthetic control (notebook 07) instead.

Notice what has *not* happened yet. No likelihood was written down, no prior was chosen, no sampler ran. The identification — parallel trends, no anticipation — did all the causal work, and four cell means did all the estimation. **This is the honest baseline against which the Bayesian section must justify itself**, and §5x will hold it to that.

The guardrail — what that confidence interval does *not* say (printed above, and it will be printed in the same words in every notebook of this cookbook): it is **not** a 90% probability that the true effect lies inside it. A confidence interval is a property of the *procedure* — 90% of intervals built this way would cover the truth across repeated samples — and this particular interval either covers the truth or it does not. The distinction is not pedantry, and here is exactly where it bites: §6 has to answer “*what is the probability the app clears its €120/store/week running cost?*” That sentence asks for a probability **about the effect itself**. The classical apparatus cannot supply one — not because its arithmetic is worse, but because that quantity does not exist in its vocabulary. It exists in a posterior. That, and not a better point estimate, is what the next section is for.

4 · Estimate — Bayesian difference-in-differences

We now put a number and an interval on that jump. One modeling choice: we first collapse each store’s 24 weekly points into a single pre-launch mean and a single post-launch mean. This is primarily the **Bertrand–Duflo–Mullainathan (2004) serial-correlation guard** — the same hazard Step 0 just demonstrated on the classical side, where it was handled by *clustering* rather than by collapsing: a store’s revenue this week is correlated with its own revenue last week (auto-correlation), so treating 24 weekly rows per store as 24 *independent* observations pretends we have far more information than we do — and that understates the standard error, making the interval look falsely sharp. Collapsing each store to a single pre mean and a single post mean throws away the within-store time series entirely, so the problem cannot arise. It is a modeling decision, not merely an API constraint — though CausalPy’s `DifferenceInDifferences` also happens to consume this classic **2×2 form** cleanly. The estimate is the `group × post` interaction coefficient β_3 , i.e. the DiD — **the same β_3 Step 0 just estimated by OLS**. Everything that follows is therefore a like-for-like comparison: same estimand, same identification, a different apparatus for expressing uncertainty about it. §5x settles the account.

The model behind `est.did`, in symbols. CausalPy 0.8.1’s `LinearRegression` fits, on design row $\mathbf{x}_i = (1, g_i, p_i, g_i p_i)^\top$ (intercept, group flag, post flag, interaction; i runs over the 80 collapsed store-periods — two rows per store):

$$\tilde{Y}_i \sim \mathcal{N}(\mathbf{x}_i^\top \beta, \sigma^2), \quad \beta_j \sim \mathcal{N}(0, 50^2), \quad \sigma \sim \text{HalfNormal}(1),$$

and those priors are **fixed by the library** (verified against the installed source), scaled for data of order one. On raw euro revenue ($\bar{Y} \approx \text{€}1,000$, residual spread in the hundreds) a HalfNormal(1) prior — a normal folded to its positive half, the natural prior for a standard deviation, which must be ≥ 0 — on σ is absurdly tight: the prior fights the likelihood and drags the effect low, with a too-narrow interval around the wrong number. So we fit on **standardized** revenue $\tilde{Y} = (Y - \bar{Y})/s_Y$ and back-transform only the interaction:

$$\widehat{\text{DiD}}_{\epsilon} = \beta_3 \cdot s_Y$$

($\bar{Y}$ is absorbed by the intercept and main effects; the coefficient of a 0/1 interaction only needs the *scale* s_Y undone). **The general lesson deserves stating as a lesson:** check your library's default prior scale against your data scale before trusting any fit — defaults tuned for standardized data will quietly mis-fit raw euros, and the failure mode is a *confident wrong answer*, not an error message.

```
DiD effect €393/store/week (true €400) * 90% credible interval (90% posterior
probability the true effect lies inside) [€274, €515]
DiD convergence: max r-hat 1.000 - min ESS 1596 - divergences 0
```


Reading the sampler's health check (DiD convergence: above). Those three numbers say whether the MCMC sampler — the algorithm that draws samples from the posterior — actually converged, and this same three-number line travels with every Bayesian fit below. **R-hat** compares the variance within each chain to the variance across the independent chains; 1.00 is perfect and ≤ 1.01 is the usual pass bar. **ESS** (effective sample size) is how many genuinely *independent* draws our autocorrelated chains are worth — a few hundred is ample for a posterior mean or interval; bigger is better. **Divergences** are steps where the sampler broke down and silently distrusts that region — **you want 0**. Under this notebook's FAST teaching profile the chains are deliberately short, so R-hat can drift to ~ 1.02 – 1.03 and PyMC may print a “*problems during sampling*” / “*effective sample size... smaller than 100*” notice above — a benign artifact of the small draw count that clears in a FULL run, not a broken model. The real backstop for trusting this estimate is the three-witnesses calibration in Step 5, not any single sampler line.

4b · Model criticism — the posterior predictive check

Before we *read an effect off* this model, we ask whether it could plausibly have **generated the data at all**. The 2×2 model is deliberately coarse — four cell means plus one shared Gaussian noise term; every store baseline α_s it doesn't describe has to live inside that noise. We draw replicated datasets from the fitted model's posterior predictive, back-transform them to euros, and overlay them on the 80 observed store-period revenues. If the real distribution (spread, lumpiness, tails) looked unlike anything the model can produce, no interval downstream would deserve trust. And notice one thing *because* it foreshadows 5c: for this coarse model to pass, its replicate spread must be **fat** — the between-store baseline scatter sits in σ — and that same fat σ is what will make the effect interval conservative.

PPC: 90% of the 80 observed store-period revenues fall inside their replicate 5-95% band ($\sim 90\%$ if calibrated) - no gross misfit.

Posterior residual sigma $\sim \text{€}170$ - note how fat: it must absorb the between-store baseline spread (sd $\text{€}150$ by construction) that the 4-mean model doesn't describe. 5c shows the price; 5d collects the refund.




5 · Validate — 2x2 cross-check and the placebo pair

Cross-check the Bayesian estimate against the 2x2 — computed here **by hand**, from four cell means, so you can see the arithmetic that Step 0's regression performs internally (the cell below verifies the two agree to the last euro). With revenue **standardized** before the fit (so CausalPy's default $O(1)$ -scaled priors don't shrink a €400 effect sitting on €1000 revenue), the two now **agree** and both land on the planted €400 with the truth inside the interval. Then **falsify**: run DiD on the *pre-period only*, splitting it into a fake “before/after” at week 6. There was no treatment then, so the placebo effect should be ≈ 0 — a large one would mean our design manufactures effects. That is placebo-in-**time**; its mirror, placebo-in-**group** (fake treated stores instead of a fake launch date), follows in 5b.

Hand-computed 2x2 €390 = Step 0's OLS interaction €390 (identical by algebra, asserted above - the regression is just bookkeeping for the four means).

Bayesian €393 * 2x2 €390 * true €400 * placebo (no real effect) €17 +/- SE €33 - within sampling noise of zero




5b · Placebo-in-group — a permutation null from the control stores

The time placebo asked “does our design manufacture effects out of **seasonality**?” The group placebo asks the complementary question: “...out of **store-to-store noise**?” Together they bracket the two ways a fake effect could arise in a panel. Take only the 20 control stores — the app never touched them, so the true effect there is *exactly zero* — declare a random half “pseudo-treated”, and run the very same 2x2 at the real launch week. Repeating over many random relabelings traces out the **permutation null**: the full distribution of DiD estimates our pipeline produces *when nothing is going on*. Two read-outs: (i) a permutation p-value — the share of relabelings whose placebo estimate beats the real one in magnitude; (ii) an **assumption-light uncertainty cross-check** — the null’s spread is how far store noise alone can move this estimator, with no likelihood, no priors, no Gaussian anywhere. File that second one away: compare it with the width of the Bayesian interval from Step 4 as you read 5c and 5d.

Permutation null: sd €14, largest |placebo DiD| €37. The real estimate €393 is ~28x the null’s sd - store noise does not produce numbers like this: permutation $p \leq 0.005$ (resolution floor 1/201).

Cross-check: noise alone moves this estimator by $\sim \pm \text{€}28$ (2 sd - and a 10-vs-10 split is even noisier than the real 20-vs-20), yet the Bayesian interval above is $\pm \text{€}121$ wide. An assumption-light hint that the 2x2 interval is conservative; 5c quantifies that, 5d fixes it.




5c · Recovery across many seeds — unbiased *and* calibrated?

A single sample can land a little off; the real test is whether the estimator is **centred on the truth over repeated samples** and whether its 90% interval **covers** the truth at the stated rate. We refit on many fresh samples (a small fast fit each; their sampler chatter is silenced below) and check both. Coverage can fail in **both** directions — an interval can be too *narrow* (overconfident) or too *wide* (conservative) — so beside counting hits we also compare the interval's width against the true seed-to-seed scatter.

DiD across 20 seeds: mean €403 (true €400) bias +3 sd €11 * 90% interval covers truth in 20/20 seeds.

But read 20/20 carefully: the headline interval's +/-€121 half-width is ~11x the true seed-to-seed scatter (sd €11) - OVER-coverage, not perfect calibration. The 2x2 collapse leaves between-store baseline spread in the residual, so the interval is honest but conservative; downstream bounds (the break-even running cost, the stress-table probabilities) inherit that width - cautious, not sharp. (Store fixed effects / within-store differencing would tighten it.)



5d · Completing the diagnosis — within-store differencing

5c’s verdict was “unbiased and covered, but roughly an order of magnitude too wide”. The cause is visible in the model spec of Step 4: the 2×2 pools stores, so the store baselines α_s (sd €150 by construction) sit in the residual σ — the PPC in 4b showed exactly that fat σ . The cure is the oldest trick in the fixed-effects book: **difference each store against itself**. Collapse every store to a single number,

$$\Delta_s = \bar{Y}_{s,\text{post}} - \bar{Y}_{s,\text{pre}},$$


which removes α_s *exactly* (it is time-invariant, so it appears in both means and cancels) while keeping the Bertrand–Duflo–Mullainathan two-period guard from Step 4. The DiD becomes a one-coefficient comparison of store-level *changes*,

$$\Delta_s \sim \mathcal{N}(\gamma_0 + \gamma_1 D_s, \sigma_\Delta^2), \quad \widehat{\text{DiD}} = \gamma_1 = \bar{\Delta}_{D=1} - \bar{\Delta}_{D=0},$$

with γ_0 the control stores’ mean change (the seasonal tide) and σ_Δ containing only averaged week-level noise, no store baselines. Same estimand, same point estimate — radically smaller residual. It’s a ten-line PyMC model with weak priors on the standardized scale. And the non-negotiable check: a *tighter* interval is progress only if coverage survives, so we re-run 5c’s seed loop on this estimator — a tightening that broke coverage would be overconfidence wearing a fix’s clothes.

Within-store DiD €391 (true €400) * 90% interval [€369, €412] - half-width €22
vs the pooled 2x2’s €121 (~6x tighter)
convergence: max r-hat 1.000 - min ESS 1962 - divergences 0
Seed loop (20 fresh rollouts): mean €403, seed-to-seed sd €9, coverage 17/20.
The 90%-interval half-width (~€17) is ~1.9x the seed scatter,
where a calibrated 90% interval sits at ~1.6x (the pooled 2x2 was at ~11x):
-> tight AND calibrated

(with only 20 seeds the hit count is itself binomially noisy).
 Decision preview: the running cost the app clears with 95% posterior probability moves from ~€274 (pooled 2x2) to ~€369/store/week (within-store posterior).



Read-out. Same point estimate, an interval several times tighter, and — the part that makes the tightening legitimate — the seed loop’s hit count (printed above) stays statistically consistent with the nominal 90% once you account for binomial noise on a handful of seeds: the width we removed was redundancy, not protection. Notice that three *independent* witnesses now agree on the honest uncertainty scale: the permutation null’s spread (5b), the seed-to-seed scatter (5c), and the within-store posterior (5d) — and a fourth, arriving from outside the Bayesian apparatus entirely: Step 0’s **cluster-robust standard error**, which landed on that same scale without a single prior (§5x lays the two side by side). Step 6 below prices the decision with the **conservative** Step-4 posterior on purpose — a rollout case that survives the wide interval is robust — and the break-even print above shows how much sharper the same case becomes once the analysis earns its precision.

5x · Point estimate vs posterior — what the Bayesian layer actually bought

SAME estimand (the ATT), SAME assumptions (parallel trends + no anticipation);
 five ways of putting an interval on it. Planted truth = €400.

ATT (€/store/week)	est	5%	95%	+/-	half
classical 2x2 * iid SE (WRONG)	390	351	430		40
classical 2x2 * clustered on store	390	370	411		20
classical within-store Delta * HC1		390	369	412	21
Bayesian 2x2 * CausalPy (Step 4)	393	274	515		121
Bayesian within-store Delta * PyMC (5d)		391	369	412	22

LOCATION - they agree. The five point estimates span just €3 (€390-€393),
 ~1% of the effect itself, and every interval covers the truth.

WIDTH - they do not. The classical CLUSTERED half-width (€20) and the within-store POSTERIOR (€22) land on the same scale, while the pooled Bayesian 2x2 is ~6x wider than either. The odd one out is the Bayesian MODEL, not the classical SE - see 4b/5c: that model keeps the store baselines in

sigma.

P(ATT > c) - a probability about the effect, at three candidate running costs:

c (€/store/wk)	classical	Bayes 2x2	Bayes within-store
120	not defined	1.00	1.00
300	not defined	0.90	1.00
380	not defined	0.57	0.80

('not defined' is not a dodge: a confidence interval has no probability to report about a hypothesis. §6's go/no-go rule needs one of these numbers.)

The honest verdict — what the posterior bought, and what it did not.

1 · They agree, and you should say so out loud. Five estimators — two classical, two Bayesian, plus the deliberately-wrong iid row — land on the same number, within a few euros of each other and of the planted €400. That is not a coincidence and it is not a disappointment: **the causal work was done by the identification, not by the machinery.** Parallel trends plus no anticipation is what turns a difference of averages into an effect; once that argument is granted, a Bayesian regression with weak priors and 960 observations is doing arithmetic the four cell means had already done. If you were hoping that “going Bayesian” would move the point estimate, this notebook is the correction: with data this plentiful and priors this weak, it must not, and it does not.

2 · The widths tell a sharper story — and it is not the flattering one. The classical clustered interval and the within-store posterior (5d) agree closely, as the print-out shows. The *pooled Bayesian 2×2* of Step 4 is the outlier — several times wider than both. And now we can name the culprit exactly, because Step 0 already did the diagnosis: that model puts the store baselines α_s (sd €150) into its residual σ , and then charges the effect for a scatter the DiD contrast never touches. The cluster-robust sandwich *refuses* to make that charge — it looks at the 40 store trajectories, sees the baselines cancel, and reports the honest width from the start. **The competently-executed classical estimator had the right interval on the first attempt; it took us until 5d to get the Bayesian model back to the same place.** Blaming the estimator would be exactly backwards: a Bayesian interval is only as good as the model behind it, and Step 4's model is coarse *by design* (it is the interface CausalPy offers). The lesson is not “Bayes is over-wide” — it is that a posterior inherits every modelling sin, silently and with a straight face, whereas the classical route forced us to make the covariance assumption explicit and therefore forced us to get it right.

3 · What the posterior *does* buy — the last block of the print-out. Look at the P(ATT > c) column. The classical row reads *not defined*, and that is the literal truth, not a rhetorical concession: the frequentist apparatus has no probability to attach to a hypothesis about a fixed parameter. Yet the €-decision in §6 is nothing but such a probability — $P(\text{lift} > \text{€120 running cost})$, then the same question under a 75% and a 50% transfer discount, then a whole stress table of them, each one propagating the full shape of the uncertainty (not just its interval endpoints) into a roll-out call. That is the payoff, and it is the only payoff:

Bayes did not rescue the estimate here — it was never in danger. Bayes bought the right to make a probability statement about the effect, which is the statement the business actually asked for.

And the last row of that little table shows why beat 2 was not an aside: at a running cost close to the effect itself, the coarse Step-4 posterior and the within-store posterior return *materially different* probabilities for the identical question, on identical data. Nothing about the world changed between

those two numbers — only the residual we were willing to charge the effect for. Width is not a cosmetic property of an interval; it *is* the euro decision once the cost gets close to the lift. §6 prices the roll-out with the conservative Step-4 posterior deliberately (a case that survives the wide interval is robust), and 5d’s print-out shows how much sharper the same case becomes once the model stops paying for scatter it never used.

Be clear about the ledger, in both directions. Bayes did **not** buy precision (the clustered SE matched it), did **not** buy a better point estimate (all five agree), and did **not** rescue us from a modelling error (it *created* one, in Step 4, that the classical SE never made). It bought the euro decision in §6, and — in notebooks with thin cells rather than 40 clean stores — it buys partial pooling and informative priors, neither of which this fat, balanced panel needed. That is the honest size of the win, and it is enough.

6 · Decide, in euros — the rollout case

Under parallel trends the DiD number is the **ATT — the effect on the pilot stores**. Rolling the app out to all 500 assumes the other 480 respond the same way, which chains rarely guarantee (pilot sites aren’t picked at random). So alongside the headline projection we **stress it two ways**: a **transfer discount** (what if the non-pilot stores realise only 75% or 50% of the pilot lift?) and a **cost sweep** (how high can the app’s running cost climb before the call flips?). The table below is P(app pays) across both — robust at full transfer, but a *pilot-further* call under aggressive discounting and higher running costs. (The €120 running cost is an illustrative all-in figure — licences + operations per store-week; one-time build/rollout capex is deliberately outside the per-week economics.)

Base case (APP_COST €120, full pilot lift across 500 stores):

```
net €273/store/week * annual value €7,097,941 [90% €3,993,808, €10,272,351]
P(app pays) 1.00 -> ROLL OUT
break-even: the app clears its cost with 95% posterior probability
while the running cost stays below ~€274/store/week.
```

P(app pays) by transfer share (rows) and running cost in €/store/week (cols):

transfer \ cost	80	120	200	300
100% (as pilot)	1.00	1.00	0.99	0.90
75%	1.00	1.00	0.96	0.46
50%	1.00	0.98	0.46	0.00

Most the roll-out can pay per store/week and still clear cost

at 90% posterior confidence:

```
at 100% (as pilot) up to ~€300
at 75%             up to ~€225
at 50%             up to ~€150
```

Read-out. At face value the base case is emphatic: $\approx$ €270/store/week of net value over the €120 running cost, $\approx$ €7M of annual value across the 500-store chain, $P(\text{app pays}) \approx 1.00 \rightarrow$ **ROLL OUT**. The stress table is where the real decision lives, and it splits along one axis: reading *across* a row, the call barely moves with the running cost while the transfer assumption holds; reading *down* a column, it degrades fast as the assumed transfer falls — at 50% transfer the €200 cost cell is roughly a coin flip (≈ 0.5) and €300 is a clear no (≈ 0.0). Since pilot stores are rarely representative

of the chain, the honest headline is the last block: the rollout can pay **up to $\approx \text{€}300/\text{store}/\text{week}$ if the whole chain responds like the pilot, but only $\approx \text{€}150$ if half the lift transfers** — that range, not the base-case $\text{€}7\text{M}$, is the number to hold against a vendor quote. If the quote lands inside it, the cheap de-risking move is a second pilot wave in deliberately *non-pilot-like* stores to measure the transfer share directly.

7 · Caveats — the assumptions under stress, live


Parallel trends is the whole ballgame and untestable post-launch; a diverging pre-trend kills the design (move to synthetic control, notebook 07). **No spillover (SUTVA)**: the app must not move revenue in *other* stores — if app-holders defect from nearby control stores, the control trend dips and DiD *overstates* the lift; if pilot buzz lifts the whole chain, it understates. No within-sample diagnostic can see either; it is a market-structure judgment (flag pilot stores that share a catchment with controls). Beyond these, two traps deserve a live demonstration each.

Trap 1 — anticipation (Assumption 2, broken on purpose). Suppose the pilot stores start behaving differently *before* week 12 — pre-launch promo buzz, staff trained early, customers holding back purchases until the app’s perks arrive — so the lift effectively starts two weeks early. The “pre” baseline is then contaminated with effect, and DiD subtracts part of the app from itself. We rerun the simulation with exactly that change and redraw Step 3’s event study:

The $k=-2$ and $k=-1$ leads sit at $\text{€}364$ and $\text{€}340$ — far outside their 90% bands (4 of 12 pre-launch leads exclude 0, vs ~ 1 expected by chance).

Naive 2x2: $\text{€}324$ vs true $\text{€}400$ — bias $\text{€}-76$, $\sim -2/12$ of the effect ($\sim \text{€}67$): the 2 contaminated pre-weeks inflate the treated baseline by that much.

Fix — drop the contaminated window from the pre-period: $\text{€}394$. (Or define 'treatment' at the ANNOUNCEMENT date rather than the launch date.)



Read-out. This is what a violated lead looks like — compare Step 3’s clean version. The last two leads jump out of their bands two weeks *before* the official launch, and because centred leads


must average zero by construction, the contamination also drags the *other* leads slightly negative: the tell is a broken, tilted pre-period, not just two high dots. The flat-leads check is no formality — it just caught a violation that quietly ate roughly a sixth of the measured effect. The remedies are cheap once you *see* it: drop the anticipation window from the pre-period (done above — the estimate snaps back to truth within noise) or date the treatment at *announcement* rather than launch. What the leads still cannot do is certify PT itself: a time-varying, group-specific shock arriving exactly at launch remains invisible.

Trap 2 — staggered adoption. Real chains rarely launch everywhere at once. With multiple adoption dates the workhorse regression becomes static **two-way fixed effects (TWFE)**:

$$Y_{st} = \alpha_s + \lambda_t + \beta^{\text{TWFE}} D_{st} + \varepsilon_{st}, \quad D_{st} = \mathbf{1}[t \geq g_s],$$

where g_s is store s 's adoption week. **Goodman-Bacon (2021):** $\hat{\beta}^{\text{TWFE}}$ is a variance-weighted average (each little 2×2 contributes a *weight* — its share of the final average; normally weights are positive and sum to one) of **every pairwise 2×2 DiD** the panel contains — early-vs-late, late-vs-early, treated-vs-untreated — *including* “forbidden comparisons” that use **already-treated** cohorts as controls. When effects grow with time since adoption, those comparisons can receive **negative weights**, so TWFE can land far from every true effect even though each one is positive. Watch it happen:

Naive TWFE €177 vs true average post-treatment effect €400 - bias -223. Use Callaway-Sant'Anna / Sun-Abraham for staggered rollouts.



Why the bias is downward here — and the cure. Cohort A adopts at week 8 and its effect keeps *growing*; when B adopts at week 16, part of TWFE’s estimate compares B against already-treated A — a “control” whose revenue is still rising with A’s growing effect. B’s lift is measured against that rising counterfactual and shrunk, and the growing tail of A’s effect looks like a calendar trend that the week effects λ_t partially absorb. Net: the largest true effects get the smallest (even negative) weights, and the estimate lands *below every single true post-adoption effect in the panel*.

The modern fix — **Callaway & Sant’Anna (2021)** — is a discipline, not a trick: *never use the already-treated as controls*. Estimate one clean 2×2 per cohort g and week t , anchored at the cohort’s last untreated week $g - 1$, using only stores **not yet treated** at t as controls:


$$\widehat{\text{ATT}}(g, t) = [\bar{Y}_{g,t} - \bar{Y}_{g,g-1}] - [\bar{Y}_{\mathcal{C}(t),t} - \bar{Y}_{\mathcal{C}(t),g-1}], \quad \mathcal{C}(t) = \{s : g_s > t\},$$

then aggregate the cells (here a simple average — cohorts are equal-sized). One **design lesson** falls straight out of the formula: in the two-cohort world above, $\mathcal{C}(t)$ is *empty* from week 16 onward — once everyone is treated there are no clean controls left and $\text{ATT}(g, t)$ simply does not exist. Real rollouts should **keep a never-treated holdout** for exactly this reason. So we regenerate the same growing-effect world with a third cohort C (20 stores, never treated) and run both estimators on it — the ~15-line loop below is the entire method.

A never-treated holdout in the panel helps TWFE (€321 here vs €177 without one, above) but does NOT fix it - the forbidden A-as-control-for-B comparisons are still inside the average.

Hand-rolled Callaway-Sant'Anna: €405 vs true average €400 - within store-noise of the truth, from 32 clean (g,t) cells; the left panel shows the cells tracking the growing effect itself.

Production tools: R ``did`` (Callaway-Sant'Anna) or ``fixest::sunab`` (Sun-Abraham); Python ``pyfixest``, ``differences``.



8 · What we tell the CMO

Finding. The loyalty app lifted pilot-store revenue by $\approx$ €400/store/week net of the seasonal tide (Step 4's posterior — and the simulation's planted truth is €400, recovered). The estimate survived every falsification we threw at it: flat pre-launch leads (Step 3), a fake launch date ($\approx$ 0, Step 5), a permutation of fake treated stores (the real estimate sits far outside anything store noise produced, 5b), a posterior predictive check (4b), and a multi-seed recovery loop (5c) — with the within-store design (5d) confirming the same point estimate under a far tighter, *calibrated* interval.

Decision. At the €120/store/week running cost the base case says **ROLL OUT** ($P(\text{app pays}) \approx 1$). But the DiD estimand is the **ATT** — **the effect on the pilot stores** — and the non-pilot stores need not respond identically, so the honest headline is the stress table's break-even band: the rollout clears its cost up to roughly **€300/store/week if the chain responds like the pilot, only $\approx$ €150 if half the lift transfers**. Hold the vendor quote against that band, not against the base-case annual € figure.

Recommended action. Phase the rollout, with wave 2 in deliberately *non-pilot-like* stores to measure the transfer share directly — and keep a small **never-treated holdout** to the very end: Step 7 showed clean controls are what keep every later wave measurable, cheap insurance an all-at-once rollout destroys.

Watch-list (the assumptions that carry everything): **parallel trends** — monitor the leads every wave; **no anticipation** — date effects from *announcement* and drop contaminated pre-weeks (the leads catch it, as demonstrated); **no cross-store spillover** — flag pilot stores sharing a catchment with controls; and **never evaluate a staggered rollout with a single post dummy** — use cohort-time effects with not-yet-treated controls (Callaway–Sant’Anna), or growing effects will bias the answer just as TWFE’s did above.

Method summary. DiD buys identification with two assumptions (PT + NA) and one subtraction; the event-study leads probe NA and make PT credible, but nothing certifies PT post-launch. Step 0’s classical read — a 2×2 OLS with store-clustered errors — reaches the same number, and (see 5x) the same honest width: the causal content is in the identification, and the Bayesian layer earns its keep at the *decision*, by turning that estimate into $P(\text{lift} > \text{cost})$. The Bayesian layer is ordinary regression — the craft is in the *scale* (default priors vs euro data, Step 4), in the *collapse* that guards against serial correlation (BDM), and in knowing which residual your interval is paying for (5b–5d). **On real data:** notebook 07b (geo-lift on a real Google geo-experiment) is this design’s nearest real-data sibling in the cookbook, and Card & Krueger (1994) remains the classic public DiD panel to practice on.