

03 · Price elasticity by region — random slopes & endogeneity

July 6, 2026

1 03 · Price elasticity by region — random slopes & endogeneity

The business decision. We set prices region by region and want each region’s **price elasticity of demand** — the standard economics measure of how sensitive demand is to price. Concretely, elasticity is the **% change in quantity sold for a 1% change in price**; it’s almost always negative (raise price -> sell less). An elasticity of -1.5 means “cut price 1%, sell 1.5% more.” We need it per region because a tourist market and a price-sensitive suburb should not be priced the same.

1.0.1 Two traps this notebook is built around

1. **Small regions are noisy.** A region with 12 weeks of data will hand you a wild elasticity estimate that’s mostly sampling noise. Acting on it is overfitting. The fix is **partial pooling** (below).
2. **Price is endogenous.** “Endogenous” means the treatment (price) is *correlated with the error term* — here because managers historically *cut prices when they expected soft demand*. So low prices coincide with low demand for reasons that have nothing to do with the price cut, which drags the naive elasticity estimate *toward zero* (makes demand look less price-sensitive than it is). No amount of clever averaging fixes this — it needs an **instrument** (notebook 11), which we preview at the end.

1.0.2 The tools, in plain terms

- **Log-log demand.** We model $\log(\text{quantity}) = \alpha + \beta \log(\text{price}) + \dots$. The magic of taking logs on both sides is that the slope β **is** the elasticity directly (a constant %-for-% relationship), so there’s one interpretable number per region.
- **Random slopes / partial pooling.** Instead of one elasticity for all regions (ignores real differences) or a separate free estimate per region (noisy), we assume each region’s elasticity β_r is drawn from a shared distribution $\beta_r \sim \mathcal{N}(\mu_\beta, \tau_\beta)$. This **hierarchical** structure makes each region’s estimate a compromise between its own noisy data and the fleet average — data-poor regions get pulled (“shrunk”) harder toward the fleet, which is exactly the protection we want. μ_β is the typical elasticity; τ_β measures how much regions genuinely differ.

Cookbook maps this to **pathmc** random slopes; here we fit the equivalent hierarchical model explicitly in PyMC so the shrinkage is fully visible. The identification logic is unchanged.

On real data. Swap in your **own transaction panel** — one row per region per week with price, units sold, and controls (competitor price, promotions, seasonality). Public analogues are retail *scanner* datasets (e.g. the Dominick’s Finer Foods store data). The

endogeneity caveat is not academic: it's the single biggest reason naive "price tests" mislead in practice.

1.1 2 · Simulate a ground truth

A region x week panel with a **constant-elasticity (log-log) demand curve** $\log Q_{rt} = \alpha_r + \beta_r \log P_{rt} + \gamma^\top Z_{rt} + \varepsilon_{rt}$, where β_r is region r 's true price elasticity (planted, varying around a fleet mean ~ -1.4). Observed confounders Z = competitor price, seasonality, trend. We start with **no** endogeneity so we can show recovery, then switch it on in step 7.

12 regions, fleet-mean elasticity -1.42; obs per region range 16-80 (regions 0-2 are thin, to exercise shrinkage)

	region	week	price	demand	competitor_price	season	\
0	region_00	2	13.003239	0.159630	20.405765	1.394170e-01	
1	region_00	3	14.245949	0.157405	16.535730	1.989368e-01	
2	region_00	12	14.310937	0.159097	20.819276	7.179470e-02	
3	region_00	13	15.626376	0.175255	21.659711	-9.648736e-17	
4	region_00	15	13.204126	0.152195	19.486540	-1.394170e-01	

	trend	log_price	log_demand
0	0.02	2.565199	-1.834900
1	0.03	2.656473	-1.848934
2	0.12	2.661024	-1.838239
3	0.13	2.748960	-1.741513
4	0.15	2.580529	-1.882592

1.2 3 · Identify — elasticity as a random slope, and the backdoor set

Estimand. Per-region elasticity $\beta_r = \partial \log Q / \partial \log P$. **Partial pooling** treats $\beta_r \sim \mathcal{N}(\mu_\beta, \tau_\beta^2)$ — noisy regions shrink toward the fleet mean μ_β , and τ_β measures how much elasticity genuinely varies. The **shrinkage factor** for region r is approximately $1 - \frac{\tau_\beta^2}{\tau_\beta^2 + \sigma^2/n_r}$ — thin regions (small n_r) get pulled harder toward the fleet. Shrinkage is a feature: it stops us over-reacting to a small region's noisy slope.

Confounding is the crux. Price is endogenous — a demand shock moves both P and Q . Conditioning on the observed backdoor set (competitor price, seasonality, trend) closes the *observable* paths. What pooling does **not** fix is an *unobserved* common shock — for that you need an instrument (notebook 11). We state this out loud rather than hide it.

1.3 4 · Estimate — three pooling regimes, side by side


We fit the same elasticity three ways to make the bias-variance trade-off concrete: **no pooling** (per-region OLS — unbiased but noisy), **complete pooling** (one fleet elasticity — stable but ignores real variation), and **partial pooling** (the hierarchical Bayes middle ground). Only partial pooling gets both the fleet level and the region-to-region spread right.

fleet-mean elasticity mu_beta -1.36 (true -1.42) · between-region spread
tau_beta 0.31

1.4 5 · Validate — recovery, shrinkage, and calibration

Estimated vs true elasticity should track the 45° line; the thin regions' point estimates get pulled toward the fleet mean (that's pooling protecting us), and the credible intervals are wider where data is thin. We also check that the 90% intervals actually **cover** the true elasticities (calibration).


MAE - no pool 0.45, complete 0.28, partial 0.21 · coverage 67%



1.5 6 · Decide, in euros — profit-maximising price per region

For a constant-elasticity demand curve with marginal cost c , the profit-maximising price is $P^* = \frac{\beta}{\beta + 1} c$ (needs $\beta < -1$). We propagate the **elasticity posterior** into a **price posterior** per region, and — the honest part — flag regions where we're not confident demand is even elastic ($\beta < -1$) for a **controlled price test** rather than acting on a shaky slope.

region	elasticity	P(elastic)	opt_price	action
region_00	-1.07	0.63	€43.4	controlled test
region_01	-1.72	1.00	€19.6	set price
region_02	-1.17	0.75	€38.4	controlled test
region_03	-1.20	0.88	€41.7	controlled test
region_04	-1.18	0.85	€45.5	controlled test
region_05	-1.29	0.97	€34.4	set price
region_06	-1.57	1.00	€22.2	set price
region_07	-1.65	1.00	€20.3	set price
region_08	-1.29	0.98	€35.1	set price
region_09	-1.17	0.85	€47.0	controlled test
region_10	-1.54	1.00	€22.7	set price
region_11	-1.50	1.00	€24.0	set price




1.6 7 · Caveats — what partial pooling does *not* fix

Partial pooling handles *noise*, not *bias*. If price is endogenous to an **unobserved** demand shock, adjusting for observed Z isn't enough and elasticity is biased **toward zero**. We demonstrate by switching the hidden shock on — and preview the IV cure.

clean fleet elasticity -1.36 (true -1.42)

endogenous fleet elasticity -0.59 (true -1.42) -> biased toward zero by +0.83

Fix = an instrument (a cost shifter) or a natural price experiment - see notebook 11 (IV).



Other caveats. Log-log assumes *constant* elasticity within region (a local approximation if elasticity varies with price level); `lag(price)` would capture stockpiling dynamics (omitted for clarity); and shrinkage is a prior choice — too few regions or a mis-specified τ_β prior can over-pool, so always inspect the fleet-vs-region spread.