

06 · Is spend *causing* sales? — causal MMM (pymc-marketing)

July 7, 2026

1 06 · Is spend *causing* sales? — causal MMM (pymc-marketing)

Runs in the legacy environment (pymc<6): `make env-legacy, kernel cmp-legacy.`

The business decision. Brand search converts like crazy — should we pour budget into it? Be careful: brand search mostly **catches** demand that upper-funnel channels (TV) and seasonality already created, so crediting it for those conversions and shifting budget its way would be a mistake.

1.0.1 The concepts

- **MMM (Marketing Mix Model)** — a regression of sales on each channel’s spend over time (usually weekly), used to attribute sales to channels and set budgets. It’s the “top-down” alternative to click-tracking.
- **Adstock (carry-over)** — advertising doesn’t only work the week it runs; its effect *decays over subsequent weeks*. Adstock transforms raw spend into an “effective” exposure that carries over.
- **Saturation (diminishing returns)** — the 10th €1,000 into a channel buys less lift than the 1st. A saturation curve bends over as spend rises; its *slope* is the **marginal** return that should drive the next euro of budget.
- **Last-click attribution** — crediting the conversion to the final touch (often brand search). It systematically **over-credits** demand-capture channels, which is the error a causal MMM fixes.
- **Incremental vs mediated** — the *incremental* contribution of a channel is what it *caused*; a channel can show a big correlation with sales while causing very little (it was downstream of the real driver).

1.0.2 Why “causal” MMM

A plain MMM regresses sales on all channels and can be fooled by a **confounder** like seasonality, which lifts sales *and* drives brand-search spend (people search more in peak season). A **causal MMM** takes a **DAG** of the market (which channels cause what) and uses the backdoor logic from notebook 05 to control for genuine confounders while *not* adjusting away legitimate mediated effects — so it separates real incremental contribution from spurious credit.

Two honest themes run through this notebook: (1) the **confounding correction is robust and decisive** — it gets the *ranking* and the go/no-go right; (2) an MMM’s **absolute** contribution levels are *not* reliable from observational data alone — so you **calibrate them with a geo experiment** (exactly Anchor B, notebook 07). Saying this out loud is the mark of an honest MMM.

On real data. Swap in your **own weekly spend-by-channel + sales** table plus known confounders (seasonality, price, distribution). Public example: the datasets shipped with Meta’s *Robyn* or *pymc-marketing*’s own tutorials. Always pair an MMM with occasional geo experiments to anchor the levels.

7-step contract.

1.1 2 · Simulate a ground truth

Weekly sales from a faithful **adstock + Michaelis-Menten saturation** process. **TV** is a big genuine driver. **seasonality** is a **confounder** — it lifts sales directly *and* drives brand-search spend. **brand_search** has only a *small* real effect, but its spend rides the seasonal wave, so a naive attribution ignoring seasonality massively over-credits it. Honest answer: TV highly incremental, brand_search only slightly.

TRUE incremental sales - TV 3,536 brand_search 333 (small)

	date_week	tv	brand_search	seasonality	sales
0	2023-01-02	42.870187	21.749830	0.000000	130.875186
1	2023-01-09	53.049712	18.348364	0.120537	136.412935
2	2023-01-16	27.623875	25.155597	0.239316	135.190674
3	2023-01-23	56.329747	32.172097	0.354605	152.617219
4	2023-01-30	43.339248	35.793793	0.464723	156.933276

1.2 3 · Identify — the marketing DAG picks controls vs mediators

DAG: **seasonality** -> {**sales**, **brand_search**}, **tv** -> **sales**, **brand_search** -> **sales**. The backdoor criterion says **seasonality is a confounder -> control for it**; omitting it is exactly the last-click mistake that over-credits brand search. The estimand for channel *c* is the interventional contribution $\mathbb{E}[\text{sales} \mid \text{do}(x_c)] - \text{baseline}$; ROI = contribution ÷ spend, each with a posterior. We hand the DAG to the MMM via `dag`, `treatment_nodes`, `outcome_node`.

Backdoor adjustment set the causal layer will use: ['seasonality']

1.3 4 · Estimate — fit the causal MMM

A caution up front, and an honest one: **MMMs are notoriously miscalibrated in absolute terms** — with channel spend correlated with seasonal demand, a short weekly series can’t fully separate “channel drove sales” from “season drove both”, so fitted contributions run *higher* than truth. That’s why practitioners **calibrate MMMs with geo experiments** (nb 07). What the MMM *does* get right, and what drives budget, is the **ranking** and each channel’s **ROI verdict**.

Output()

Causal MMM contribution - TV 11,775 brand_search 2,083
(true: TV 3,536, brand_search 333 - absolute levels run high; see note.)


1.4 5 · Validate — confounding correction, ROI with uncertainty, and saturation

Three views:

1. **The confounding correction** (robust): naive attribution *omitting* seasonality over-credits brand_search; adjusting for it (what the DAG prescribes) collapses that credit toward the truth.
2. **ROI with uncertainty** — a forest of each channel's posterior ROI vs the break-even line; this is the decision-relevant quantity, and it's far more stable than the absolute contribution.
3. **Saturation curves** — contribution vs spend per channel, showing diminishing returns (the shape that determines *marginal* ROI and hence where the next euro should go).

brand_search credit - naive 4,843 | adjusted 679 | true 333

ROI - TV 2.81x [90% 0.29,8.02] · brand_search 0.82x [90% 0.24,0.89]



How to read this. *Left* is the headline confounding correction, and it's robust: the **naive** attribution (orange) that ignores seasonality massively over-credits brand_search, while **adjusting for seasonality** (gold) collapses that credit toward the small truth — the causal MMM (blue) does the same. *Middle* — the decision-relevant quantity, each channel's **ROI with a credible interval**: TV's interval sits above the break-even line (1x) and brand_search's sits below it, which is the whole budget story in one picture. *Right* — the **saturation curves**: contribution flattens as spend rises (diminishing returns), so the *next* euro into a channel is worth its curve's *slope*, not its average — which is why we reallocate in small steps rather than dumping the whole budget into the highest-average-ROI channel. The one honest caveat (Step 7): these ROI *levels* are approximate and should be anchored with a geo experiment; the *ranking* is what's trustworthy.

1.5 6 · Decide, in euros — reallocate budget on marginal ROI

Budget chases **incremental ROI**, not last-click credit. Even with approximate absolute levels, the **ROI ordering is decisive**: brand_search sits below break-even (mostly capturing seasonal demand), TV clears it. We estimate the sales impact of shifting a slice of budget from the low-ROI to the high-ROI channel — the concrete action — and flag that the *levels* need a geo-test calibration.

Shifting 15% of brand_search spend (€383) to TV -> projected +€764 sales

P(TV ROI > brand_search ROI) = 0.96 -> reallocate toward TV

Caveat: this uses AVERAGE ROI; with saturation, marginal ROI falls as you pour more into TV -

so reallocate in steps and re-estimate, and calibrate the absolute levels with a geo experiment.

1.6 7 · Caveats

- **Absolute contributions need calibration.** Here they ran ~2-3x the truth because channel spend rides the seasonal wave the MMM can't fully net out from a short series. **Trust the ranking and the ROI verdict, not the absolute euros** — calibrate levels with a geo experiment (nb 07). This is the single most important thing to say about any MMM.
- **The DAG is the assumption.** If TV genuinely creates incremental search value, encode that edge or you'll under-credit it.
- **Adstock/saturation priors matter.** Carry-over (λ) and diminishing-returns (α , κ) are weakly identified from limited weekly data; use informative priors from experiments where you have them.
- **Correlated channels -> wide, correlated ROI posteriors.** When two channels move together, the data can't fully separate them; the intervals will (honestly) show it.